

LaMachine: A meta-distribution for NLP software

Maarten van Gompel and Iris Hendrickx

Centre for Language and Speech Technology (CLST)

Radboud University, Nijmegen, the Netherlands

`proycon@anaproj.nl`, `i.hendrickx@let.ru.nl`

`https://proycon.github.io/LaMachine`

Abstract

We introduce LaMachine, a unified Natural Language Processing (NLP) open-source software distribution to facilitate the installation and deployment of a large amount of software projects that have been developed in the scope of the CLARIN-NL project and its current successor CLARIAH. Special attention is paid to encouragement of good software development practices and reuse of established infrastructure in the scientific and open-source software development community. We explain what LaMachine is, how it can be used and the technical details. We also compare LaMachine to alternative software distributions and discuss its advantages and limitations. We illustrate how LaMachine can be used in two case studies, one in an exploratory text mining project at the Dutch Health Inspectorate where LaMachine was applied to create a research environment for automatic text analysis for health care quality monitoring, and a second case where LaMachine was used to create a workspace for a one-week, intense collaboration by a diverse research team.

1 Introduction

Software is a key deliverable and a vital component for research in projects such as those under the CLARIN umbrella. Software can be seen as an invaluable instrument that enable researchers to perform their studies. It is CLARIN's core mission to make digital language resources, including software, available to the wider research community.

We see that NLP software often takes on complex forms such as processing pipelines invoking various individual components, which in turn rely on various dependencies. Add dedicated web-interfaces on top of that and you obtain a suite of interconnected software that is often non-trivial to install, configure, and deploy. This is where LaMachine comes in.

Software is a broad but well-ingrained notion in society nowadays, referring to any form of computer program. This can manifest in various forms, whether it is an app on a phone, a graphical desktop application, a command line tool, or a web-based application. These different *interfaces* generally address different *audiences*; the data scientist will feel at home at the command line and in scripting environments and uses these fairly low-level interfaces, whereas researchers in the humanities demand higher-level graphical interfaces. There is often a power trade-off between lower and higher-level interfaces, with the former providing maximum flexibility at the cost of a steeper learning curve, technical ability, and a do-it-yourself mentality. Higher level interfaces, on the other hand, expose certain functionality of the software in an easy and accessible way, but in doing so often can not expose the full power of the software. The cost trade-off is also apparent in the construction of the interfaces, where high-level interfaces are typically far more costly to build.

LaMachine incorporates software providing different types of interfaces which, as seen above, typically address different audiences. Whilst we attempt to accommodate both technical¹ and less-technical audiences², there is a natural bias towards the former as lower-level interfaces are often a prerequisite to build higher-level interfaces on. Depending on the *flavour* (more on this later) of LaMachine chosen, it

¹Data scientists, DevOps, system administrators, developers.

²The wider researcher community, particularly the Humanities; also educational settings.

makes a good virtual research environment for a data scientist, whether on a personal computer or on a computing cluster, a good development environment for a developer or a good deployment method for production servers in for example CLARIN centres.

This paper provides a detailed description of LaMachine, its purpose and functionality and the technical details. We discuss the advantages and limitations of LaMachine compared to alternative software distributions in Section 6. We demonstrate how LaMachine can create a fully functioning and standalone research environment for text mining and NLP for Dutch texts in two use cases in Section 7.

2 What is LaMachine?

LaMachine is an open-source NLP software distribution. LaMachine facilitates the installation, distribution and configuration of software. It does not fork, modify or appropriate the participating software in any way, nor does it provide a hosting place or repository for software. We classify LaMachine as a *meta distribution* as it can be installed in various contexts. The heart of LaMachine consists of a set of machine-parsable instructions on how to obtain, build (e.g. compile from source), install and configure software. This is notably different from the more classical notion of Linux distributions, which generally provide their own repositories with (often binary) software packages. LaMachine builds on this already established infrastructure by taking these repositories as a foundation where possible. Similarly, as implied in point five above, there are different programming-language-specific ecosystems providing their own repositories, such as the Python Package Index³ for Python, CRAN⁴ for R, CPAN⁵ for Perl, Maven Central⁶ for Java. LaMachine relies on those to obtain and install software. In doing so, we compel participating software projects to adhere to well-established distribution standards and ensure the software is more sustainable (van Gompel et al., 2016). Moreover, we ensure that LaMachine never becomes a prerequisite for the software but merely a courtesy or convenience, and attempt to limit any amount of duplication in packaging and distribution efforts.

LaMachine lives in an open-source ecosystem and therefore builds on Unix-like platforms; this primarily means Linux, as well as BSD and, with some restrictions, macOS. This by definition excludes certain software for different platforms, such as mobile platforms (Android/iOS/etc), native Windows/mac desktop applications, or certain interface types in general such as classical desktop GUI applications or mobile ‘apps’, all of which fall beyond our scope. Cygwin⁷ is not tested or supported either. However, virtualisation technology enables deployment on a wider range of platforms, including Windows.

LaMachine exists since May 2015 and has been installed extensively ever since by numerous users. Quantifying its actual use is not trivial despite having some tracking solutions in place, but our statistics show that LaMachine was installed at least over a thousand times, on machines in over 25 different countries since its inception⁸. In early 2018 version 2 was released which was a significant redesign, powered by Ansible⁹. We ourselves use LaMachine as a means by to make our full software stack available to end-users, whether those be individual researchers, research groups, or hosting providers like CLARIN centres. LaMachine has in fact been adopted in all of those areas. We also see a regular influx of questions on our issue tracker, indicating community interest.

The software included in LaMachine has to adhere to the following prerequisites:

1. The software must have a recognised (i.e. OSI-approved) open-source license. Proprietary (closed-source) software is explicitly excluded.
2. The source code must be hosted in a public version controlled repository.¹⁰
3. There must be a build process to compile the source, if applicable, and install the program or library.

³<https://pypi.org>

⁴<https://cran.r-project.org/>

⁵<https://www.cpan.org>

⁶<https://search.maven.org>

⁷A unix environment on Windows

⁸<https://applejack.science.ru.nl/lamastats/lamachinestats.html>

⁹<https://www.ansible.com>

¹⁰e.g. Github, Gitlab, Bitbucket, provided the repository is public.

4. There must be some kind of release protocol (adhering to semantic versioning) that publishes software using the proper technology-specific channels. We will elaborate on this later.
5. All software that is incorporated in LaMachine must bear at least some relevance to the field of Natural Language Processing.
6. Participating software must be actively maintained (i.e. not outdated or abandoned) and not place any demands on outdated dependencies.

3 Included Software

In LaMachine, software is grouped into various “packages”, each package¹¹ groups one or multiple programs that have some kind of relation. The installation manifest lists all packages that will be installed, at the user’s discretion. After the initial installation, the user can always add more packages using `lamachine-add` or editing the installation manifest directly.

LaMachine was initially conceived as the primary means of distribution of the software stack developed at the Language Machines Research Group and the Centre of Language and Speech Technology, Radboud University Nijmegen. The majority of this software was either fully or partially developed under the auspices of CLARIN-NL or successor CLARIAH. Some software by other CLARIN-NL/CLARIAH partners is also included. LaMachine is not limited to one research group and is explicitly open to participation by other software providers, especially those also in CLARIAH.

We list a selection of the most important software included in LaMachine, grouped by research institute:

- by the Language Machines Research Group and the Centre of Language and Speech Technology, Radboud University, Nijmegen:¹²
 - **Timbl** – A memory-based machine learning toolkit, and **Mbt**, a memory-based tagger based on timbl. Python bindings included as well.
 - **Ucto** – A multilingual rule-based tokeniser. Python binding included as well.
 - **Frog** – An integration of various memory-based natural language processing (NLP) modules developed for Dutch. It can do Part-of-Speech tagging, lemmatisation, named entity recognition, shallow parsing, dependency parsing and morphological analysis. Also included in LaMachine; Python bindings for Frog and **Toad**, Trainer Of All Data, training tools for Frog.
 - **Wopr** – Memory-based Word Predictor.
 - **CLAM** – Quickly build RESTful webservices, powers many webservices offered by LaMachine.
 - **FoLiA** – Format for Linguistic Annotation (van Gompel and Reynaert, 2013), with tools and libraries in/for Python and C++.
 - **FLAT** – FoLiA Linguistic Annotation Tool: a web-based linguistic annotation tool.
 - **PyNLPL** – Python Natural Language Processing Library.
 - **Colibri Core** – Colibri core is an NLP tool as well as a C++ and Python library for working with basic linguistic constructions such as n-grams and skipgrams (i.e patterns with one or more gaps, either of fixed or dynamic size) in a quick and memory-efficient way.
 - **Gecco** – Generic Environment for Context-Aware Correction of Orthography, an NLP pipeline for spelling correction, and **Valkuil.net**, an instantiation thereof for Dutch.
 - **PICCL** – A set of workflows (NLP pipeline) for corpus building through OCR, post-correction (through **TICCL**) and Natural Language Processing.

¹¹For those familiar with Ansible, a package in LaMachine is an Ansible role

¹²Links: <https://languagemachines.github.io/timbl>, <https://languagemachines.github.io/mbt>, <https://languagemachines.github.io/ucto>, <https://languagemachines.github.io/frog>, <http://ilk.uvt.nl/wopr>, <https://proycon.github.io/folia>, <https://proycon.github.io/clam>, <https://github.com/proycon/flat>, <https://proycon.anaproy.nl/pynlpl>, <https://proycon.github.io/colibri-core/>, <https://github.com/proycon/gecco>, <https://github.com/proycon/valkuil-gecco>, <https://github.com/LanguageMachines/PICCL>, <https://github.com/LanguageMachines/ticcltools>, <https://github.com/proycon/labirinto>, <https://github.com/proycon/oersetter-webservice>

- **Labirinto** – A web-based portal listing all available tools in LaMachine, an ideal starting point for LaMachine.
- **Oersetter** – A Frisian-Dutch Machine Translation system in collaboration with the Fryske Akademy.
- by the University of Groningen: **Alpino**¹³ – A dependency parser and tagger for Dutch.
- by the Vrije Universiteit Amsterdam: **KafNafParserPy**¹⁴ – a python module to parse NAF files, another format for linguistic annotation; and **NafFoLiApy**¹⁵ – converters between FoLiA and NAF.
- by Utrecht University: **T-scan**¹⁶ – a Dutch text analytics tool for readability prediction (initially developed at TiCC, Tilburg University, and in collaboration with Radboud University, Nijmegen)
- by the Meertens Instituut: **Python Course for the Humanities**¹⁷ – interactive tutorial and introduction into programming with Python for the humanities.

In addition to the above listed specific software, LaMachine also incorporates a large number of renowned tools by external international parties, offering most notably a mature Python environment with scientific modules. The following list gives an impression and is not exhaustive:¹⁸

- **Python:** Numpy, Scipy, Matplotlib, Scikit-Learn, IPython, Jupyter, ...
 - **Jupyter Lab** The successor of the popular Jupyter Notebooks, offers notebooks, a web-based IDE, terminals. An ideal entry point to get started with LaMachine and all it contains!
 - **PyTorch** - Deep-learning library for Python
 - **NLTK** - Natural Language Toolkit for Python
 - **Spacy** - Industrial-Strength NLP in Python
- **R, Perl, Java:**
 - **NextFlow** - A system and language for writing parallel and scalable pipelines in a portable manner.
 - **Stanford CoreNLP** - Various types of linguistic enrichment
- **Tesseract** - Open Source Optical Character Recognition (OCR)
- **Tensorflow** - Open-source machine learning framework
- **Kaldi** - Speech Recognition Framework (ASR)
- **Moses** - Statistical Machine Translation system

4 Architecture

In this section we present the technical design choices that were made and lay out the architecture of LaMachine.

4.1 Flavours

LaMachine provides ample flexibility that allows it to be deployable in different contexts. First of all there is flexibility regarding the installation form, which we call *flavours*:

1. **Local installation** – This installs the bulk of LaMachine in a separate local virtual environment (a separate directory¹⁹) that has to be explicitly activated to be used. This also allows for multiple different installations of LaMachine on the same host system (e.g. for different users or with different software configurations). This local installation still actively relies on various global dependencies that are available through the package manager of your distribution.

¹³<http://www.let.rug.nl/vannoord/alp/Alpino/>

¹⁴<https://github.com/cltl/KafNafParserPy>

¹⁵<https://github.com/cltl/NafFoLiApy>

¹⁶<https://github.com/proycon/tscan>

¹⁷<http://www.karsdorp.io/python-course/>

¹⁸**Links:** <https://jupyterlab.readthedocs.io/en/stable/>, <https://pytorch.org>, <http://www.nltk.org>, <https://spacy.io>, <http://www.nextflow.io>, <https://stanfordnlp.github.io/CoreNLP/>, <https://github.com/tesseract-ocr/tesseract>, <https://tensorflow.org>, <http://kaldi-asr.org>, <http://www.statmt.org/moses>

¹⁹Python users should know we just use `virtualenv` for this, with some additions of our own

2. **Global installation** – This flavour is used for a host that is fully dedicated to LaMachine. Everything will be installed globally so there is only one installation possible, multiple users will all make use of the same installation. Unlike the local installation, we do not make use of a Python Virtual Environment in this flavour.
3. **Docker container** – Installs LaMachine in a container. Containerisation separates the entire run-time environment from the host system, only the Linux kernel is shared. It is a lighter option than full virtualisation²⁰. Support for two other forms of containerisation, i.e. Singularity²¹ and LXC/LXD²², are currently under development.
4. **Virtual Machine** — Installs LaMachine as a virtual machine, i.e. through full virtualisation. This allows deployment of LaMachine on hosts which would otherwise not support it, such as Windows. Virtualisation for LaMachine is achieved through Vagrant and VirtualBox²³.
5. **Remote provisioning** – Installs LaMachine on a remote dedicated server. This option is most suited for hosting centres and directly uses Ansible’s remote provisioning abilities.

The different flavours all offer a different degree of separation from the host OS, where Virtual Machines are completely virtualised, containers still share the kernel with the host OS, and the two native installation flavours, local and global, actually compile against the machine’s distribution itself and thus offer the least amount of overhead.

The two native options support a variety of major GNU/Linux distributions: Debian, Ubuntu, Arch Linux, CentOS, Fedora and, to a more limited degree, we also support macOS, powered in part by the homebrew package manager²⁴. Support for macOS is limited because not all participating software supports it natively²⁵. Certain Linux Distributions that are derivatives of the aforementioned distributions *may* also work, such as RedHat Enterprise Linux (CentOS) and Linux Mint (Ubuntu).

For the containerisation and virtualisation solutions, the default distribution we supply is Debian²⁶. It is, however, still possible to build your own container or virtual machine based on any of the other supported distributions. In fact, containers, virtual machines and remote provisioning can all be considered special wrapped forms of the global installation flavour.

Pre-built docker containers and virtual machine images with a limited selection of participating software are uploaded to the Docker Hub²⁷ and Vagrant Cloud²⁸, respectively, for each LaMachine release.

4.2 Versions

LaMachine offers three distinct *versions*, regardless of the flavour:

1. **stable** - This is the default and recommended for most situations. It installs the latest stable versions of all included software.
2. **development** - This installs the latest development versions of the included software. In practise, this usually means that software is pulled directly from the version-controlled repository and is compiled and installed from source. Due to the experimental nature, the development version of LaMachine may at times break and not install successfully.
3. **custom** - This installs custom versions of all included software, i.e. the user explicitly specifies which versions to install for each software package. Such a version list can for instance be exported from another LaMachine installation, and then allows to rebuild a similar environment from scratch, providing a limited level of reproducibility. We say limited, because packages provided by the underlying distribution are not a part of this scheme.

²⁰Note that Docker on other platforms such as Windows and macOS does use full virtualisation.

²¹<https://www.sylabs.io/singularity/>

²²<https://linuxcontainers.org>

²³<https://vagrant.org>, <https://www.virtualbox.org>

²⁴<https://brew.sh>; alternatives such as macports are not supported

²⁵One can always still fall back to the virtualised flavour of LaMachine

²⁶The latest stable release

²⁷<https://hub.docker.com/r/proycon/lamachine/>

²⁸<https://app.vagrantup.com/proycon/boxes/lamachine>

The nomenclature is admittedly a bit confusing, but the notion of version discussed in this section refers to the versions of the various software packages inside LaMachine. It should not be confused with the actual *version release number* of LaMachine as a whole, which is the version number assigned to the collective of installation scripts LaMachine provides, and which marks LaMachine releases.

4.3 Bootstrapping

Installation of LaMachine begins with a single *bootstrap* command²⁹ executed on the command line. It can interactively query the user for her software preferences (*stored as the host configuration*), e.g. the flavour of LaMachine, as well as the set of software to install, *the installation manifest*. This set is never static but can be customised by the user. The bootstrap procedure, a screenshot of which is shown in Figure 1, detects and installs the necessary prerequisites automatically and eventually invokes Ansible to perform the bulk of the work, unless a pre-published container or VM is selected. Figure 2 provides a schematic view of the LaMachine architecture.

```

proycon@mhysa ~ $ bash <(curl -s https://raw.githubusercontent.com/proycon/LaMachine/master/bootstrap.sh)
=====
LaMachine v2.3.1 - NLP Software distribution
~) (http://proycon.github.io/LaMachine)
(----i Language Machines research group
// \ Centre of Language and Speech Technology
/ / // Radboud University Nijmegen
=====

Welcome to the LaMachine Installation tool, we will ask some questions how
you want your LaMachine to be installed and guide you towards the installation
of any software that is needed to complete this installation.

Where do you want to install LaMachine?
1) in a local user environment
   installs as much as possible in a separate directory
   for a particular user; can exist alongside existing
   installations. May also be used (limited) by multiple
   users/groups if file permissions allow it. Can work without
   root but only if all global dependencies are already satisfied.
   (uses virtualenv)
2) in a Virtual Machine
   complete separation from the host OS
   (uses Vagrant and VirtualBox)
3) in a Docker container
   (uses Docker and Ansible)
4) Globally on this machine
   dedicates the entire machine to LaMachine and
   modifies the existing system and may
   interact with existing packages. Usually requires root.
5) On a remote server
   modifies an existing remote system! Usually requires root.
   (uses ansible)
Your choice? [12345] 1
Where do you want to create the local user environment?
By default, a directory will be created under your current location, which is /home/proycon
If this is what you want, just press ENTER,
Otherwise, type a new existing path:
Where do you want to create the local user environment? [press ENTER for /home/proycon]
LaMachine comes in several versions:
1) a stable version; you get the latest releases deemed stable (recommended)
2) a development version; you get the very latest development versions for testing, this may not always work
3) custom version; you decide explicitly what exact versions you want (for reproducibility);
   this expects you to provide a LaMachine version file (customversions.yml) with exact version numbers.
Which version do you want to install? [123] 1

```

Figure 1: A screenshot of the bootstrap procedure

Once LaMachine is installed, in any of its flavours, it can be updated from inside by running `lamachine-update`. This updates all of the software managed by LaMachine.

4.4 Metadata

LaMachine aims to harmonise the metadata of all installed software. This is accomplished by converting metadata from upstream repositories, i.e. the repositories where tool providers deposit their software, to a common yet simple standard called CodeMeta³⁰ (Jones et al., 2016; Boettiger, 2017) where possible, or encouraging software developers to provide their codemeta metadata inside their source code repositories and using that directly. CodeMeta is a linked data initiative that provides a mapping from/to various commonly used software metadata standards³¹. All this metadata in LaMachine in turn enables other

²⁹See <https://proycon.github.io/LaMachine>

³⁰<https://codemeta.github.io/>, described in JSON-LD

³¹such as DOAP, Github API, Debian packages, Python distutils, R packages, Ruby gems, Maven metadata, DataCite, Wiki-Data

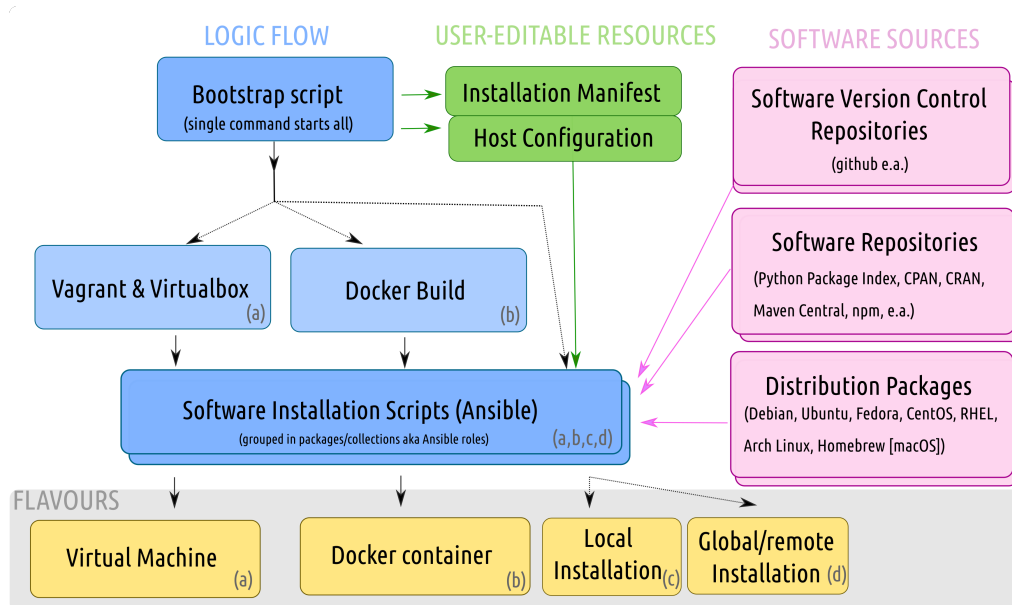


Figure 2: A schematic representation of the LaMachine architecture

tools to do proper service discovery and provenance logging.

The software metadata discussed here is of a generic and practical nature, it most importantly contains information regarding the version, source, licensing, and authors of the software. The underlying principle is to obtain software metadata from as close to the source as possible. This can be contrasted with other metadata efforts, also within CLARIN, to establish a more centralised metadata store³² with more specialised and research-oriented metadata. There is merit to both approaches as they serve distinct functions. But in order to best serve everybody’s different metadata needs, these bottom-up and top-down processes need to come to some kind of synthesis, e.g. by having component registries do automatic harvesting of CodeMeta (or other) metadata, similar to what LaMachine does.

5 Interfaces and Audiences

We already addressed the need for different interfaces for different audiences in Section 1. The challenge we face, and for which LaMachine offers a solution, is one of software maintainability, distribution and deployment. How do we maintain, distribute, and deploy software that is often highly complex and consists of multiple interconnected components, considering that this software is used differently by different audiences? A key aspect here is the reusability and accessibility of individual software components. The philosophy we subscribe to encourages the development and distribution of software in a layered or modular fashion, allowing each building block to serve as the foundation for another more high-level interface, without sacrificing the usability of the foundation as such. This is in line with the UNIX philosophy of developing tools that do *“one thing only, and do it well”* and is contrasted with monolithic software solutions that are limiting because they either provide only high-level interfaces but do not expose their inner components to build upon, or they are as a swiss-army knife full of needless but inseparable components.

5.1 Low-level Interfaces

In any of its flavours, LaMachine offers low-level shell access, i.e. a command-line interface accessed through a terminal. In flavours that are separated from the host system by a network, this is accomplished over `ssh`. After accessing the LaMachine environment through a terminal, as shown in Figure 3, the user has the liberty to do whatever she wants and can invoke any of the tools that offer a text interface, including text editors such as `vim`, `emacs` or `nano`, version control systems such as `git` and interpreters such as `Python` or `R` (or compilers for that matter). This allow users to use one of the many specialised

³²Such as CLARIN’s CMDI Component Registry; <https://www.clarin.eu/content/component-metadata>

programming libraries included in LaMachine to build their own tools and . All this makes LaMachine ideally suited as a development environment.

```

proycon@mhyza ~ $ source lamachine-activate
=====
,           LaMachine v2.4.0 - Unified NLP Software distribution
(----)      (https://proycon.github.io/LaMachine)
/| |\       Centre for Language and Speech Technology
/ / ||      Radboud University Nijmegen
=====

Build Name: dev
Version: development      Build time: 2018-08-31 12:13:49 CEST
Maintainer: proycon@mhyza.anaproy.nl

Welcome to LaMachine!
- run lamachine-list to see a verbose list of all installed software
  add -s for a short sorted view
- run lamachine-add to add extra software collections to this installation
  add the --list flag to see a list of installable packages
- run lamachine-update to update and test your LaMachine installation
  add the --edit flag to edit settings and/or the installation manifest directly
- run lamachine-start-webserver to (re)start the webserver

proycon@mhyza ~ (dev) $

```

Figure 3: Activation of a local LaMachine environment on the command-line

5.2 High-level Interfaces: Web-based access

LaMachine comes with a webserver³³. This webserver serves various web-capable tools that are incorporated in LaMachine. One of these tools is a portal website that provides an overview of all installed software, and acts as a point of access to all its web services and web applications. This portal, as shown in Figure 4, leverages the metadata registry compiled in each LaMachine installation (see section 4.4).

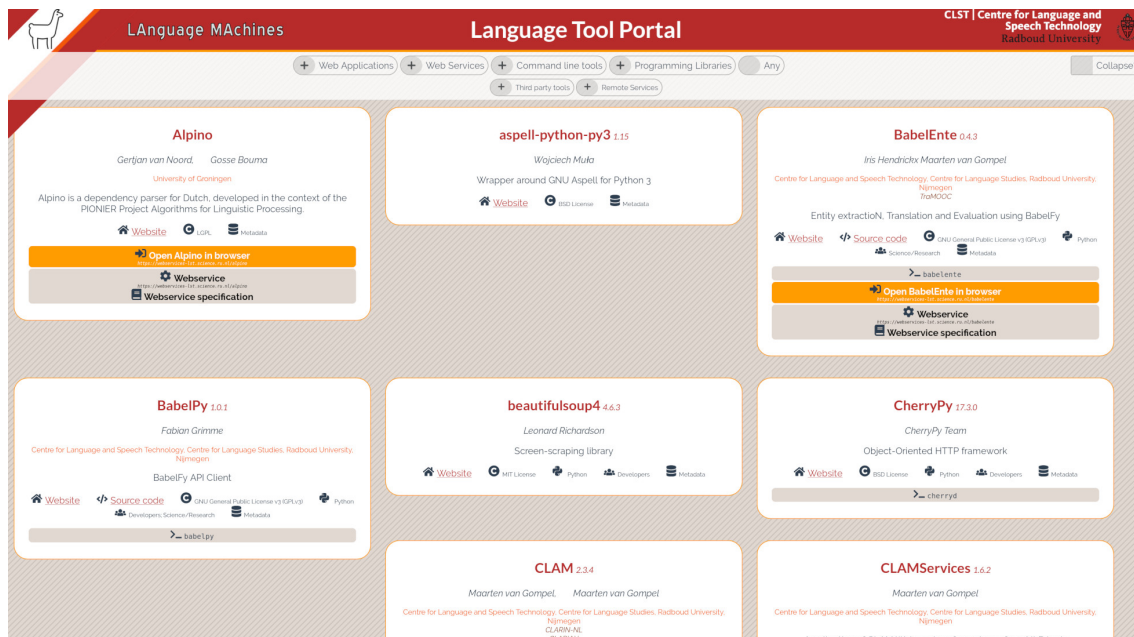


Figure 4: A screenshot of the LaMachine portal, powered by Labirinto: <https://github.com/proycon/labirinto>

LaMachine also comes with a Jupyter Lab³⁴ installation which provides a web-based Integrated Development Environment (IDE) for scripting in Python and R, web-based terminal access, and so-called *notebooks* which mix text, code and data output and have gained great popularity in the data science community. This is shown in Figure 5. This type of virtual laboratory provides a powerful interface that

³³In situations where web interfaces are not needed or desired, the user has the ability to opt-out of this.

³⁴<https://jupyter.org/>

provides an alternative to regular shell access and may have a larger appeal for those parts of the audience that do not feel completely comfortable in the terminal.

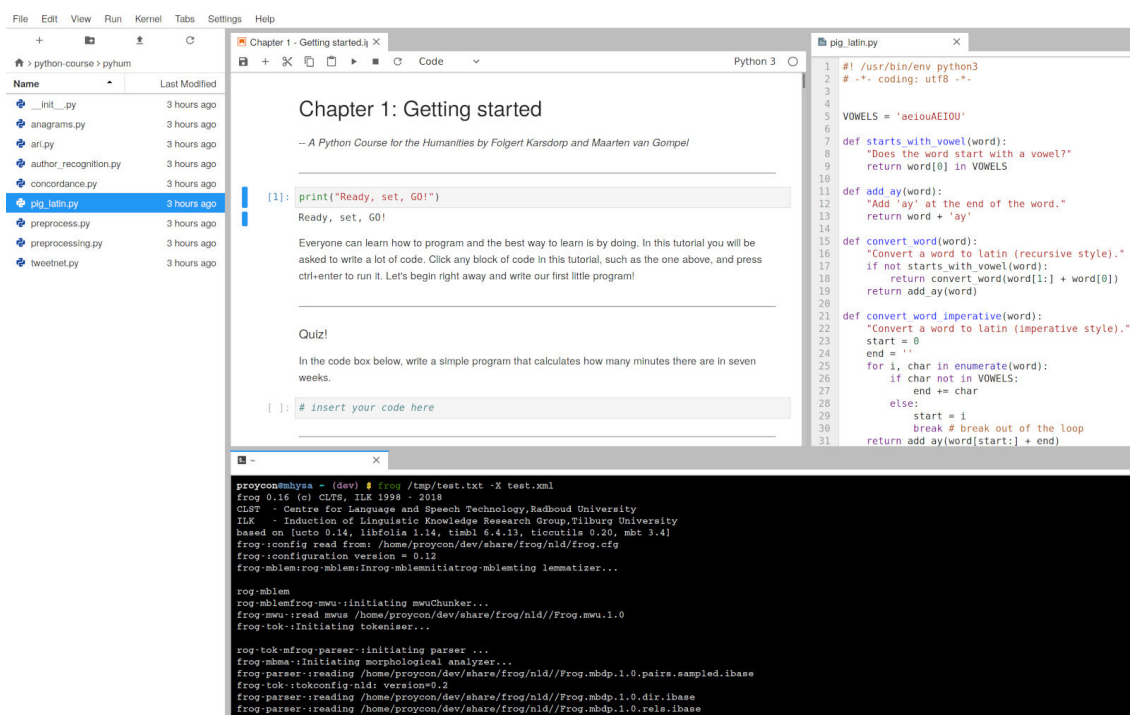


Figure 5: A screenshot of the Jupyter Lab Environment

LaMachine provides various CLAM-based webservices. CLAM is a long-time CLARIN project that allows developers to turn their command-line NLP tool into a RESTful webservice with relatively little effort (van Gompel and Reynaert, 2014), and which in addition also provides a generic web-interface for human end-users. The portal website in LaMachine links to these. This again addresses an audience that does not require a high degree of technical expertise, but is suitable for a wider audience of researchers, including those in the Humanities.

6 Comparison to alternative distributions

How does LaMachine compare to alternatives? To answer this question we must first look at what may be considered alternatives to LaMachine. At this point, it is important to emphasise once more that LaMachine is offered as convenience. LaMachine always seeks to use installation channels that the user can also access directly. For instance, is the user only interested in a single Python package that LaMachine offers? Then she can just as well install it directly using `pip install`. Is the user only interested in a single package that the underlying Linux distribution (or macOS's homebrew) already provides? Then, by all means, she should go for that.

LaMachine starts to fill a need when the user is interested in complex software, such as complex NLP pipelines with multiple components, the individual installation (sometimes involving compilation from source) and configuration of which would not be trivial. Moreover, LaMachine shines when the user wants higher-level interfaces on top of that, something she can use right out of the box. It also fulfills a role in case the user simply does not know what she wants or needs yet, and needs an environment where she can explore a variety of pre-installed tools and experiment.

Within the scope of CLARIN, one might be inclined to draw comparisons with the CLARIN Language Resource Switchboard (Zinn, 2016b), which also offers access to NLP tools through a web interface where users can upload their texts for NLP processing, providing easy access and automated tool discovery. This, however, is notably different from what LaMachine is. The CLARIN switchboard is not a software distribution, it links to externally hosted applications; and LaMachine does currently not provide any switchboard functionality, but merely offers a portal interface. Moreover, LaMachine is used

by data scientists and researchers who do have programming experience and want to work with complex NLP software, on their own machines and unencumbered by high-level interfaces if they so desire.

There are also CLARIN initiatives that provide not just a switchboard, but a more elaborate high-level virtual research environment (VRE) with graphical interfaces to accommodate the researcher and allow execution and visualisation of various NLP pipeline. The Dutch CLARIAH WP3 project³⁵ is such an initiative and LaMachine is currently developed in close collaboration with this project, in a role of technology supplier.

Comparisons could be drawn between LaMachine as a kind of Linux distribution and other Linux distributions. That too, however, would not be appropriate as we characterise LaMachine more as a meta-distribution which simply builds on the foundation of several existing major Linux distributions. As a meta-distribution that runs in a variety of flavours and on a variety of platforms, it offers a great deal of flexibility that normal distributions do not have. At the same time, the scope of LaMachine is more narrowly defined than that of a generic Linux distribution, including even specialised Linux distributions that focus on scientific computing, such as Fedora Scientific³⁶.

The best comparison can perhaps be drawn with Anaconda³⁷, which we could also qualify as a meta-distribution under our definition, and which enjoys great popularity in the data science community. Anaconda has a focus on Python and R. It is much wider in scope than LaMachine, which has a strong NLP and CLARIN/CLARIAH focus. When it comes to Python, which Anaconda was initially aimed at, there is considerable overlap in the modules that LaMachine and Anaconda offer. Unlike LaMachine, Anaconda does not focus on providing different flavours such as a VM or a Docker Container³⁸, nor separate stable and development versions of the included software.

At the start of the development cycle for LaMachine v2, we investigated whether Anaconda and its ecosystem would provide a sufficient foundation for us to build upon. We concluded this was not the case for a number of reasons: Anaconda introduces more overhead than we desired for our purposes and conflicted with certain technology choices we made. We wanted our native flavours (local environment and global environment) to be as close to the underlying distribution as possible and to reuse existing technologies and packages, such as compilers (e.g. gcc) and interpreters (e.g. python), from the distribution (or from homebrew for macOS). Anaconda, on the other hand, chooses to provide its own packages and builds on those. It does so using its own package manager (conda) and package repositories (conda-forge). This would require repackaging certain software for the anaconda ecosystem.

For the distribution and deployment of complex software setups such as NLP pipelines, containers are a common solution nowadays. The Docker flavour of LaMachine provides something similar, but rather than providing a single static Docker recipe that builds a single kind of container, LaMachine offers a high degree of flexibility for the construction of different containers, i.e. containing different software or having different configurations, based on the user's needs. It is quite feasible to instantiate a variety of LaMachine containers and use a container orchestration system such as Kubernetes or Docker Swarm for automated deployment, scaling and management thereof. Such orchestration, however, is beyond the scope of LaMachine itself, which aims to provide a singular environment, i.e. everything installed in a LaMachine instance shares the same userspace and duplication is strongly minimised.

The flexibility LaMachine offers, with various flavours and versions, makes it more accessible and deployable in multiple contexts, but this does come at a cost. Maintaining LaMachine itself is a non-trivial task that requires constant maintenance and testing³⁹; software exists in an ever moving ecosystem. Any of the distributions we target, or software that participates, *might* at any time introduce a change that requires us to adapt LaMachine accordingly. Similarly, underlying Linux distribution releases that are currently up-to-date and supported, will eventually be deprecated and replaced by newer releases, a development LaMachine is obliged to follow by design.

³⁵<https://github.com/meertensinstituut/clariah-wp3-vre>

³⁶<https://labs.fedoraproject.org/en/scientific/>

³⁷<https://www.anaconda.com/>

³⁸This is not a technical limitation but simply a matter of different objectives, one could easily create a VM or a container and install Anaconda in it.

³⁹We run automated tests on a continuous integration platform to this end

7 Case studies

LaMachine is most often used by a researcher or data scientist who installs LaMachine on his own computer (in a local environment, Docker container or VM). Here we discuss two case studies where LaMachine was used under circumstances that go beyond this most common usage of LaMachine: using LaMachine on a secure server without internet connection, and using LaMachine to provide a work environment for a group of researchers. These case studies illustrate the advantages and limitations of LaMachine in practice.

7.1 Text Mining for Health Inspection

We participated in a small Dutch national project titled “*Text mining for Inspection: an exploratory study on automatic analysis of health care complaints*”⁴⁰ led by IQhealthcare⁴¹, the scientific centre for healthcare quality of RadboudUMC hospital. This project took place at the Dutch Health Inspectorate and aimed to apply text mining techniques to health care complaints that have been registered at the national contact point for health care (Landelijk Meldpunt Zorg⁴²). We investigated the usefulness of text mining to categorise and cluster complaints, to automatically determine the severity of incoming complaints, to extract patterns and to identify risk cases. This project turned out to be a good test case of the applicability and usefulness of LaMachine as a standalone research environment. As the complaint data is highly sensitive, it could not leave the secure servers of the health inspectorate and was stored in an environment without internet access. We needed to bring the software to the data via a shared folder.

We used a virtual machine (VM) image of LaMachine and we ran this 64-bits Linux-based VM inside another VM with Windows Server 2012, provided to us by the health inspectorate for this project, in which we did have administrative rights but no internet access. In terms of hardware we ran on a machine with 8 cores and 32GB internal memory available. Note that as we had no internet access available on the working server, we prepared a custom-made VM image of LaMachine that included an additional editor and some extra Python libraries. LaMachine provided a fully functional research environment and we ran all our experiments within LaMachine. We interacted with LaMachine both through the command line, which offers a standard shell and enables access to all lower-level tools and programming languages; and through the (offline) webbrowser to use the Jupyter Notebook environment.

LaMachine comes with some simple data sharing facilities that allowed us to access the sensitive complaint data via a single shared dataspace between host and the VM. Extensive data search and management functions are deliberately beyond the scope of LaMachine, and left to more high-level tooling.

We used many of the available tools in LaMachine within this project: Frog for linguistic annotation of the textual content of the complaint and the scikit-learn Python package for classification, T-scan for feature extraction in the form of text characteristics and colibri-core for n-gram analysis.

7.2 Workshop: Cataloguing of Textual Cultural Heritage Objects

The ICT-Research Platform Netherlands and NWO organise a yearly one-week workshop ‘ICT with Industry’⁴³ to stimulate collaboration between industry and academia. The industrial partner provides a problem and a team of researchers from different backgrounds and universities collaborate to come up with solutions. We participated in the 2019 edition on the case study by the Dutch Royal Library who wanted to investigate automatic methods for cataloguing of textual cultural heritage objects, in this particular case a large collection of digital dissertations.

For this workshop, computing power was purchased at SURFsara, a collaborative organization for ICT in Dutch education and research. Subsequently, we had the ability to create Virtual Machines on their hosting platform. For this workshop we had twelve participants and decided to create a single multi-core and high memory VM to share amongst the participants to create one common digital workspace.

⁴⁰<https://bit.ly/2N2AICS>

⁴¹<http://www.iqhealthcare.nl/nl/>

⁴²<https://www.landelijkmeldpuntzorg.nl>

⁴³<https://ict-research.nl/ict-with-industry/ictwi2019/>

As recommended by SURFsara, we opted for one of their default Linux images as a basis for the VM, based on Ubuntu 18.04, instead of providing a LaMachine VM image directly like we did in 7. This was recommended because their image was already preconfigured to integrate nicely with their cloud environment. Inside this VM we simply bootstrapped the local installation flavour of LaMachine. Here we benefit from the flexibility LaMachine offers because of its various flavours, and regardless of the flavour, the resulting installations are always functionally equivalent. The local environment flavour had the added bonus of not being in anyone's way in case any of the twelve workshop participants did not want to use LaMachine, considering it needs to be explicitly activated prior to usage.

LaMachine offered a convenient platform for a range of different explorations and experiments in the area of NLP and text mining. However, for some situations LaMachine, or rather Linux in general, was not a good fit for the audience of the workshop: for team members who did not have experience with a non-Windows environment, LaMachine was not a suitable or useful tool. The limit of LaMachine was also reached for members who wanted to use desktop text editors with a graphical user interface as this is not offered by LaMachine. Moreover, we did not manage to get X-forwarding working in the Ubuntu Linux VM and after a few attempts the team gave up on resolving this issue due to time pressure. This, also demonstrates that fine-tuning the configuration of certain aspects of LaMachine, but especially beyond LaMachine, is beyond the reach of a data scientist without system administration skills. This certainly also applies also to the installation as a whole in the SURFsara context, which involved things like the partitioning, formatting and mounting of (virtual) drives and setting up user accounts on the shared VM, all of which require some system administration skills and are too context-specific to be within the scope of LaMachine. LaMachine was convenient and speeded up writing code as the most common scientific data-related packages are already present in LaMachine.

8 Conclusion & Future work

LaMachine provides a flexible solution for the installation of a variety of NLP software, resulting in a kind of virtual laboratory that, through various interfaces, can be employed by a variety of people, from developers and data scientists to the wider research community that CLARIN explicitly addresses.

Most users of LaMachine are installing the software distribution in their own machine, either locally, via Docker or VM. However, we described two case studies in section 7 where we explore the possibilities under more constrained, complex, and demanding circumstances, which contributes to thorough testing and running of LaMachine.

Aside from the incorporation of new relevant software, the main objectives for the future are to provide greater *interoperability* between the included tools through better *high-level interfaces* for the researcher. We see this as a bottom-up process and have now established a firm foundation to build upon. Note that such proposed interfaces, including the current portal application in LaMachine, are always considered separate independent software projects, which may be deployed by/in/for LaMachine, but also in other contexts. LaMachine remains 'just' a software distribution at heart.

An important part of our future focus will be on interoperability with the higher-level tools emerging from the CLARIAH WP3 VRE effort, but also with other parts of the CLARIN infrastructure; single-sign on authentication being a notable example here.

LaMachine v2 is open to outside contribution. Contributor documentation has been written, and at this stage, we greatly welcome external participants to join in.

Acknowledgements

This research was funded by NWO CLARIN-NL, CLARIAH and the ZonMw project *Tekstmining in het toezicht: een exploratieve studie naar de automatische verwerking van klachten ingediend bij het Landelijk Meldpunt Zorg*, project number 516004614. We thank the Dutch Health Inspectorate, IQhealthcare, and Tim Voets for their valuable contributions and help in the ZonMw project, and NWO, the ICT Research Platform Nederland, the Lorentz Centre, the Royal Library and all our team members for making the ICT with Industry workshop a success.

References

- [Boettiger2017] C. Boettiger. 2017. Generating CodeMeta Metadata for R packages. *The Journal of Open Source Software*, 2:454.
- [Jones et al.2016] MB. Jones, C. Boettiger, A. Cabunoc Mayes, A. Smith, P. Slaughter, K. Niemeyer, Y. Gil, M. Fenger, K. Nowak, M. Hahnel, et al. 2016. CodeMeta: an exchange schema for software metadata. *KNB Data Repository*.
- [van Gompel and Reynaert2013] M. van Gompel and M. Reynaert. 2013. FoLiA: A practical XML format for linguistic annotation - a descriptive and comparative study. *Computational Linguistics in the Netherlands Journal*, 3.
- [van Gompel and Reynaert2014] M. van Gompel and M. Reynaert. 2014. CLAM: Quickly deploy nlp command-line tools on the web. In *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: System Demonstrations*, pages 71–75. Dublin City University and Association for Computational Linguistics.
- [van Gompel et al.2016] M. van Gompel, J. Noordzij, R. de Valk, and A. Scharnhorst. 2016. Guidelines for Software Quality. CLARIAH Task 54.100.
- [Zinn2016a] C. Zinn. 2016a. The CLARIN Language Resource Switchboard. *Proceedings of the CLARIN Annual Conference. CLARIN ERIC*.
- [Zinn2016b] Claus Zinn. 2016b. The clarin language resource switchboard. In *Proceedings of the CLARIN Annual Conference, Aix-en-Provence, France*.