LaMachine: A meta-distribution for NLP software

Maarten van Gompel
Centre for Language and Speech Technology (CLST)
Radboud University, Nijmegen, the Netherlands
`proycon@anaproy.nl`
`https://proycon.github.io/LaMachine`

Abstract

We introduce LaMachine, a unified Natural Language Processing (NLP) open-source software distribution to facilitate the installation and deployment of a large amount of software projects that have been developed in the scope of the CLARIN-NL project and its current successor CLARIAH. Special attention is paid to encouragement of good software development practises and reuse of established infrastructure in the scientific and open-source software development community.

1 Introduction

Software is a key deliverable and a vital component for research in projects such as those under the CLARIN umbrella. It is software that provides researchers the instruments to yield for their research; without it a lot of research would become highly unfeasible or right-out impossible. It is CLARIN's core mission to make digital language resources, including software, available to the wider research community.

We see that NLP software often takes on complex forms such as processing pipelines invoking various individual components, which in turn rely on various dependencies. Add dedicated web-interfaces on top of that and you obtain a suite of interconnected software that is often non-trivial to install, configure, and deploy. This is where LaMachine comes in.

LaMachine incorporates software providing different types of interfaces[1] that typically address different audiences. Whilst we attempt to accommodate both technical[2] and less-technical audiences[3], there is a natural bias towards the former as lower-level interfaces are often a prerequisite to build higher-level interfaces on. Depending on the *flavour* of LaMachine chosen, it makes a good virtual research environment for a data scientist, whether on a personal computer or on a computing

---

[1] Command line interfaces, programming interfaces, web-user interfaces, webservices.

[2] Data scientists, DevOps, system administrators, developers.

[3] The wider researcher community, particularly the Humanities; also educational settings.

cluster, a good development environment for a developer or a good deployment method for production servers in for example CLARIN centres.

2   Architecture

Being an open-source NLP software distribution, LaMachine is constrained to POSIX-compliant platforms; this primarily means Linux, but also BSD and, with some restrictions, macOS. This stands in contrast with mobile platforms (Android/iOS/etc), native Windows/mac desktop software, or certain interface types in general such as classical desktop GUI applications, all of which fall beyond the scope of LaMachine. Software that is incorporated must 1) bear some relevance to NLP, 2) be under a recognised open-source license, 3) be deposited in a public version controlled repository[4] and 4) have a release protocol (with semantic versioning) using the proper technology-specific channels.

LaMachine is a *meta distribution* as it can be installed in various contexts. At its core, LaMachine consists of a set of machine-parsable instructions on how to obtain, install and configure software. These are implemented using Ansible[5]. This is notably different from the more classical notion of Linux distributions, which generally provide their own repositories with (often binary) software packages. LaMachine builds on this already available infrastructure by taking these repositories as a given and only needs to know which repositories to use. Similarly, there are different programming-language-specific ecosystems providing their own repositories, such as the Python Package Index[6] for Python, CRAN[7] for R, CPAN[8] for Perl, Maven Central[9] for Java. LaMachine again relies on those to pull and install software from and never forks, archives, or modifies the software in any way. In doing so, we compel participating software projects to adhere to well-established distribution standards and ensure the software is more sustainable towards the future [van Gompel et al.2016]. Moreover, we ensure that LaMachine never becomes a prerequisite for the software but merely a courtesy or convenience.

LaMachine provides ample flexibility that allows it to be deployable in different contexts. First of all there is flexibility with regard to the target platform, where we support several major GNU/Linux distributions (Debian, Ubuntu, CentOS, RedHat Enterprise Linux, Fedora, Arch Linux), as well as macOS. Second, there is flexibility with regard to the form, where we support *containerisation* through Docker[10], *virtualisation* through Vagrant[11] and VirtualBox[12], direct re-

---

[4]e.g. Github, Gitlab, Bitbucket

[5]https://www.ansible.com

[6]https://pypi.org

[7]https://cran.r-project.org/

[8]https://www.cpan.org

[9]https://search.maven.org

[10]https://www.docker.com

[11]https://vagrant.org

[12]https://www.virtualbox.org

mote provisioning through Ansible (for production servers), or an installation that is either global to the machine or local in a custom directory for a specific user. Pre-built docker containers and virtual machines with a limited selection of participating software are regularly uploaded to the Docker Hub and Vagrant Cloud, respectively.

Installation of LaMachine begins with a single bootstrap command[13]. It may interactively query the users for their software preferences, e.g. the flavour of LaMachine and the set of software to install, which is never static but can be customized by the user. The user may also opt for installing the latest releases, the more experimental development versions of the software, or specific custom versions (to facilitate scientific reproducibility). The bootstrap procedure detects and installs the necessary prerequisites automatically and eventually invokes Ansible to perform the bulk of the work.

After a successful build, the user may interact with LaMachine either through the command line, which offers a standard shell and enables access to all lower-level tools and programming languages; or through his or her webbrowser, which when enabled offers a simple portal page towards all installed web-capable tools. Amongst these services is also a Jupyter Notebook environment[14], which offers a more graphical scripting environment for e.g. Python and R that has gained great popularity in the scientific community.

The current version of LaMachine comes with some simple data sharing facilities only: we provide a single shared dataspace between host and VM/container (where applicable), or allow simple upload. Extensive data search and management functions are deliberately beyond the scope of LaMachine, and left to more high-level tooling.

With LaMachine we also attempt to harmonise the metadata of all installed software, by converting metadata from upstream repositories to the CodeMeta standard[15] [Jones et al.2016, Boettiger2017] where possible, or encouraging software developers to provide their own. This in turn enables other tools to do proper service discovery and provenance logging.

3 Software

As LaMachine was conceived initially[16] as the primary means of distribution of the software stack developed at CLST, Radboud University Nijmegen, it includes a lot of our software. A full list goes beyond the scope of this overview; we will merely mention some CLARIN-NL/CLARIAH-funded tools: ucto (a tokeniser), Frog (an NLP suite for Dutch), FoLiA (Format for Linguistic Annotation, with assorted tools), FLAT (a web-based linguistic Annotation tool), PICCL (an OCR

---

[13]See `https://proycon.github.io/LaMachine`

[14]`https://jupyter.org/`

[15]`https://codemeta.github.io/`, described in JSON-LD

[16]LaMachine exists since May 2015 and has been used extensively ever since by numerous users

and post-OCR correction pipeline) and CLAM. We already include some relevant software by the University of Groningen and VU Amsterdam as well. Moreover, LaMachine incorporates a large number of renowned tools by external international parties, offering most notably a mature Python environment[17], but also R and Java.

4   Conclusion & Future work

The recent release of LaMachine v2, which constituted a full rewrite, has opened up LaMachine to outside contribution. Contributor documentation has been written, and at this stage, we greatly welcome external participants to join in. Aside from the incorporation of new relevant software, the main objectives for the future are to provide greater *interoperability* between the included tools through better *high-level interfaces* for the researcher. We see this as a bottom-up process and have now established a firm foundation to build upon. Note that such proposed interfaces, including the current portal application in LaMachine, are always considered separate independent software projects, which may be deployed by/in/for LaMachine, but also in other contexts. LaMachine remains 'just' a software distribution at heart.

Development of LaMachine presently takes place in the scope of the CLARIAH WP3 Virtual Research Environment (VRE) project[18], which has higher ambitions in accommodating the researcher and connectivity of data and services, and transcends also those of the CLARIN Language Resource Switchboard [Zinn2016]. An important part of our future focus will therefore be on interoperability with the higher-level tools emerging from the VRE efforts, but also with other parts of the CLARIN infrastructure; single-sign on authentication being a notable example here.

References

[Boettiger2017] C. Boettiger. 2017. Generating CodeMeta Metadata for R packages. *The Journal of Open Source Software*, 2:454, November.

[Jones et al.2016] MB. Jones, C. Boettiger, A. Cabunoc Mayes, A. Smith, P. Slaughter, K. Niemeyer, Y. Gil, M. Fenner, K. Nowak, M. Hahnel, et al. 2016. CodeMeta: an exchange schema for software metadata. *KNB Data Repository*.

[van Gompel et al.2016] M. van Gompel, J. Noordzij, R. de Valk, and A. Scharnhorst. 2016. Guidelines for Software Quality. CLARIAH Task 54.100.

[Zinn2016] C. Zinn. 2016. The CLARIN Language Resource Switchboard. *Proceedings of the CLARIN Annual Conference. CLARIN ERIC*.

---

[17]scipy, numpy, scikit-learn, matplotlib, nltk, spacy, pytorch, keras, gensim, etc...

[18]https://github.com/meertensinstituut/clariah-wp3-vre